

Predictions Based on 2011-2023 Home Attendance Records

This document seeks to utilize attendance records of Duke University home football games from the previous 12 seasons (2011-2023) to predict the number of attendees at Duke football home games during the 2024 season.

Packages

```
library(tidyverse)
library(tidymodels)
```

Importing the Dataset

```
att_data <- read_csv("data/Duke Stats - DukeAttendanceV3.csv")

att_data <- att_data |>
  mutate(Day = as.factor(Day))

home_att_data <- att_data |>
  filter(Site == "Home", Year < 2024)

home_opp_list <- c("Elon", "Connecticut", "Florida St.",
                  "North Carolina", "SMU", "Virginia Tech")

home_att_data
```

```
# A tibble: 84 x 51
  OppName      OppFPI DukeFPI FPI_diff DukeFPI_NetChange OppFPI_PrevYear
  <chr>      <dbl>   <dbl>   <dbl>         <dbl>         <dbl>
1 Richmond      NA     -6.1     NA           -2.1           NA
2 Stanford     24.4    -6.1    30.5          -2.1          24.2
3 Tulane     -20.3    -6.1   -14.2          -2.1         -17.3
4 Florida St.  15.3    -6.1    21.4          -2.1          17.2
5 Wake Forest  -0.2    -6.1     5.9          -2.1           -6
6 Virginia Tech 11.8    -6.1    17.9          -2.1          18.4
7 Georgia Tech   5       -6.1    11.1          -2.1           5.3
8 Florida Int'l  -8      -1.7    -6.3           4.4          -5.1
9 N.C. Central  NA      -1.7     NA           4.4           NA
10 Memphis    -13.2   -1.7   -11.5           4.4         -24.6
# i 74 more rows
# i 45 more variables: FPI_Diff_PrevYear <dbl>, Surface <chr>, Month <dbl>,
#   Date <dbl>, Year <dbl>, Day <fct>, Start_Time <dbl>, Site <chr>,
#   Result <chr>, DukePts <dbl>, OppPts <dbl>, PointDiff <dbl>, AttNum <dbl>,
#   AttPct <dbl>, ESPN_WinPred <dbl>, COVID_Limit <lgl>, Rain <lgl>,
#   City <chr>, State <chr>, TV_Coverage <chr>, Bowl <lgl>,
#   DukeRankGametime <dbl>, OppRankGametime <dbl>, OppRankSeasonEnd <dbl>, ...
```

Attendance History for 2024 Opponents

Wallace Wade Stadium capacity:

- Pre-rennovation: 33,941 (1982–2015)
- Post-rennovation: 40,004 (2016-present)

All Teams

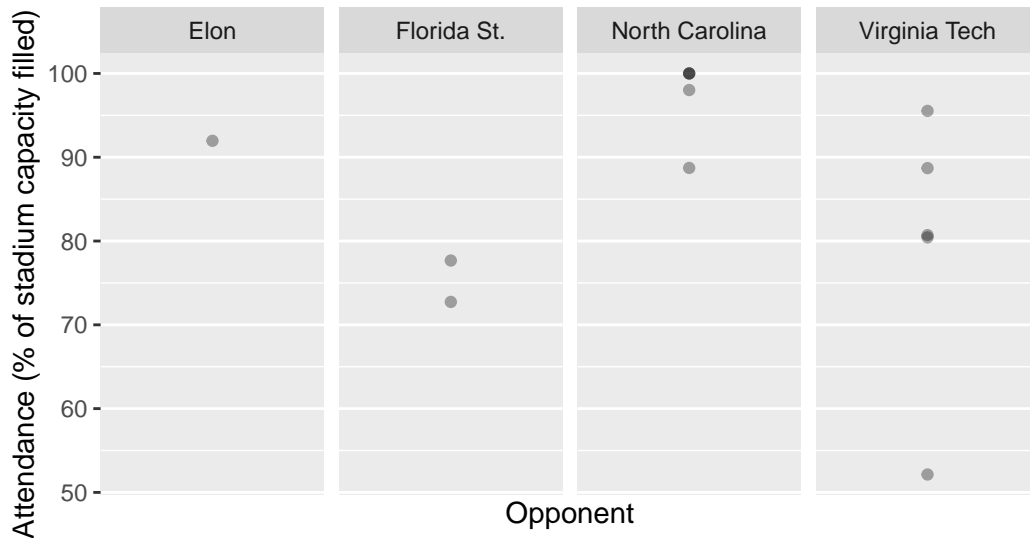
```
home_att_data |>
  filter(OppName %in% home_opp_list) |>
  ggplot(
    aes(x = 0, y = AttPct)
  ) +
  geom_point(alpha = 0.333) +
  facet_wrap(~OppName, strip.position = "top", nrow = 1) +
  scale_x_continuous(labels = NULL, breaks = NULL) +
  labs(title = "Duke Home-Game Attendance per Opponent",
```

```

subtitle = "Percentage of Stadium Capacity Filled per Game\n2011-2023",
x = "Opponent",
y = "Attendance (% of stadium capacity filled)"

```

Duke Home-Game Attendance per Opponent
Percentage of Stadium Capacity Filled per Game
2011-2023




```

home_att_data |>
  filter(OppName %in% home_opp_list) |>
  ggplot(
    aes(x = 0, y = AttNum)
  ) +
  geom_point(alpha = 0.4) +
  facet_wrap(~OppName, strip.position = "top", nrow = 1) +
  scale_x_continuous(labels = NULL, breaks = NULL) +
  labs(title = "Duke Home-Game Attendance per Opponent",
       subtitle = "Number of Attendees per Game\n2011-2023",
       x = "Opponent",
       y = "Attendance (# of people)")

```

Duke Home-Game Attendance per Opponent

Number of Attendees per Game
2011–2023



Elon

```
home_att_data |>
  filter(OppName == "Elon") |>
  summarize("Name" = OppName,
            "End-of-Season FPI" = OppFPI,
            Month,
            Date,
            Year,
            "# of Attendees" = AttNum,
            "% of Stadium Capacity Filled" = AttPct)
```

A tibble: 1 x 7

```
  Name `End-of-Season FPI` Month Date Year `# of Attendees`
  <chr>          <dbl> <dbl> <dbl> <dbl>          <dbl>
1 Elon              NA     8    30  2014          31213
# i 1 more variable: `% of Stadium Capacity Filled` <dbl>
```

Connecticut

```
home_att_data |>
  filter(OppName == "Connecticut") |>
  summarize("Name" = OppName,
            "End-of-Season FPI" = OppFPI,
            Month,
            Date,
            Year,
            "# of Attendees" = AttNum,
            "% of Stadium Capacity Filled" = AttPct)
```

```
# A tibble: 0 x 7
# i 7 variables: Name <chr>, End-of-Season FPI <dbl>, Month <dbl>, Date <dbl>,
#   Year <dbl>, # of Attendees <dbl>, % of Stadium Capacity Filled <dbl>
```

UConn never faced against Duke in Wallace Wade Stadium from 2011 to 2023.

Florida St.

```
home_att_data |>
  filter(OppName == "Florida St.") |>
  summarize("Name" = OppName,
            "End-of-Season FPI" = OppFPI,
            Month,
            Date,
            Year,
            "# of Attendees" = AttNum,
            "% of Stadium Capacity Filled" = AttPct)
```

```
# A tibble: 2 x 7
  Name      `End-of-Season FPI` Month  Date  Year `# of Attendees`
  <chr>                <dbl> <dbl> <dbl> <dbl>          <dbl>
1 Florida St.             15.3    10    15  2011           24687
2 Florida St.             13.3    10    14  2017           31073
# i 1 more variable: `% of Stadium Capacity Filled` <dbl>
```

North Carolina

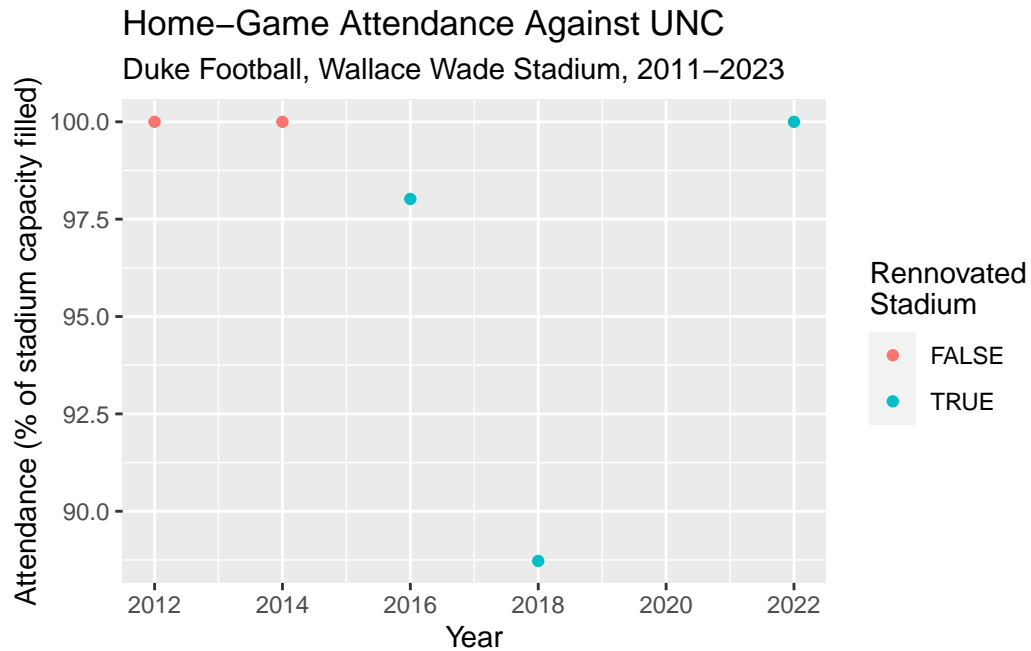
```
home_att_data |>
  filter(OppName == "North Carolina") |>
  summarize("Name" = OppName,
            "End-of-Season FPI" = OppFPI,
            Month,
            Date,
            Year,
            "# of Attendees" = AttNum,
            "% of Stadium Capacity Filled" = AttPct)
```

A tibble: 6 x 7


	Name	`End-of-Season FPI`	Month	Date	Year	`# of Attendees`
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	North Carolina	10.6	10	20	2012	33941
2	North Carolina	4.4	11	20	2014	33941
3	North Carolina	14	11	10	2016	39212
4	North Carolina	-2.6	11	10	2018	35493
5	North Carolina	10.2	11	7	2020	NA
6	North Carolina	6.2	10	15	2022	40004

i 1 more variable: `% of Stadium Capacity Filled` <dbl>

```
home_att_data |>
  filter(OppName == "North Carolina") |>
  ggplot(
    aes(x = Year, y = AttPct, color = Rennovated)
  ) +
  geom_point() +
  scale_x_continuous(breaks = seq(from = 2012, to = 2023, by = 2)) +
  labs(title = "Home-Game Attendance Against UNC",
       subtitle = "Duke Football, Wallace Wade Stadium, 2011-2023",
       x = "Year",
       y = "Attendance (% of stadium capacity filled)",
       color = "Rennovated\nStadium")
```



```
home_att_data |>
  filter(OppName == "North Carolina") |>
  ggplot(
    aes(x = Year, y = AttNum, color = Renovated)
  ) +
  geom_point() +
  scale_x_continuous(breaks = seq(from = 2012, to = 2023, by = 2)) +
  labs(title = "Home-Game Attendance Against UNC",
        subtitle = "Duke Football, Wallace Wade Stadium, 2011–2023",
        x = "Year",
        y = "Attendance (# of attendees)",
        color = "Renovated\nStadium")
```



SMU

```
home_att_data |>
  filter(OppName == "SMU") |>
  summarize("Name" = OppName,
            "End-of-Season FPI" = OppFPI,
            Month,
            Date,
            Year,
            "# of Attendees" = AttNum,
            "% of Stadium Capacity Filled" = AttPct)
```

```
# A tibble: 0 x 7
# i 7 variables: Name <chr>, End-of-Season FPI <dbl>, Month <dbl>, Date <dbl>,
#   Year <dbl>, # of Attendees <dbl>, % of Stadium Capacity Filled <dbl>
```

UConn never faced against Duke in Wallace Wade Stadium from 2011 to 2023.

Virginia Tech


```
home_att_data |>
  filter(OppName == "Virginia Tech") |>
  summarize("Name" = OppName,
            "End-of-Season FPI" = OppFPI,
            Month,
            Date,
            Year,
            "# of Attendees" = AttNum,
            "% of Stadium Capacity" = AttPct)
```

A tibble: 6 x 7

	Name	End-of-Season FPI	Month	Date	Year	# of Attendees
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	Virginia Tech	11.8	10	29	2011	27392
2	Virginia Tech	7.9	11	15	2014	30107
3	Virginia Tech	13.7	11	5	2016	38217
4	Virginia Tech	3.4	9	29	2018	32177
5	Virginia Tech	7.3	10	3	2020	NA
6	Virginia Tech	-6.2	11	12	2022	20857


i 1 more variable: % of Stadium Capacity <dbl>

```
home_att_data |>
  filter(OppName == "Virginia Tech") |>
  ggplot(
    aes(x = Year, y = AttPct, color = Rennovated)
  ) +
  geom_point() +
  scale_x_continuous(breaks = seq(from = 2011, to = 2023, by = 2)) +
  labs(title = "Home-Game Attendance Against Virginia Tech",
       subtitle = "Duke Football, Wallace Wade Stadium, 2011-2023",
       x = "Year",
       y = "Attendance (% of stadium capacity filled)",
       color = "Rennovated\nStadium")
```

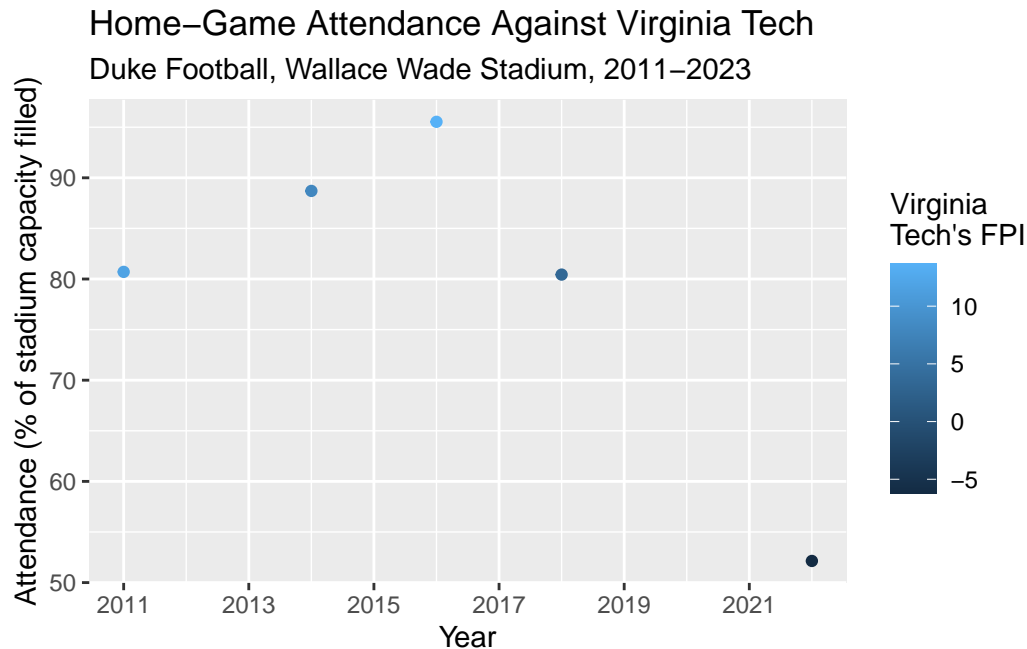


```
home_att_data |>
  filter(OppName == "Virginia Tech") |>
  ggplot(
    aes(x = Year, y = AttNum, color = Renovated)
  ) +
  geom_point() +
  scale_x_continuous(breaks = seq(from = 2011, to = 2023, by = 2)) +
  labs(title = "Home-Game Attendance Against Virginia Tech",
        subtitle = "Duke Football, Wallace Wade Stadium, 2011-2023",
        x = "Year",
        y = "Attendance (# of attendees)",
        color = "Renovated\nStadium")
```

Home-Game Attendance Against Virginia Tech Duke Football, Wallace Wade Stadium, 2011–2023



```
home_att_data |>
  filter(OppName == "Virginia Tech") |>
  ggplot(
    aes(x = Year, y = AttPct, color = OppFPI)
  ) +
  geom_point() +
  scale_x_continuous(breaks = seq(from = 2011, to = 2023, by = 2)) +
  labs(title = "Home-Game Attendance Against Virginia Tech",
        subtitle = "Duke Football, Wallace Wade Stadium, 2011-2023",
        x = "Year",
        y = "Attendance (% of stadium capacity filled)",
        color = "Virginia\nTech's FPI")
```



Team Performance vs. Attendance

Can football team performance – both of Duke and its opponent – be used to predict the attendance turnout of future Duke home games?

Previous-Season FPI

This section will seek to determine if the Football Power Index (FPI) of an opposing team at the end of one season is a decent predictor of home-game audience turnout in the *following* season.

```
home_att_data_prevFPI <- home_att_data |>
  filter(!is.na(OppFPI_PrevYear)) |>
  mutate(OppFPI_PrevYear = OppFPI_PrevYear,
         FPI_Diff_PrevYear = FPI_Diff_PrevYear)

home_att_data_prevFPI |>
  ggplot(
    aes(x = OppFPI_PrevYear, y = AttPct, color = FPI_Diff_PrevYear)
  ) +
```

```
geom_point() +
geom_smooth(method = "lm") +
labs(title = "Previous-Season Opponent FPI vs. Current-Season Attendance,\n
per Season",
      subtitle = "Duke University Football, Wallace Wade Stadium, 2011-2023",
      color = "FPI Difference\nbetween Duke\nand opponent",
      x = "Opponent FPI (at end of previous season)",
      y = "Game Attendance (% of stadium capacity)")
```


`geom_smooth()` using formula = 'y ~ x'

Warning: Removed 5 rows containing non-finite values (`stat_smooth()`).

Warning: The following aesthetics were dropped during statistical transformation: colour
 i This can happen when ggplot fails to infer the correct grouping structure in the data.

i Did you forget to specify a `group` aesthetic or to convert a numerical variable into a factor?

Warning: Removed 5 rows containing missing values (`geom_point()`).



```
prev_fpi_glm <- linear_reg() |>
  set_engine("lm") |>
  fit(AttPct ~ OppFPI_PrevYear, data = home_att_data_prevFPI)

tidy(prev_fpi_glm)
```

```
# A tibble: 2 x 5
  term          estimate std.error statistic  p.value
<chr>         <dbl>    <dbl>    <dbl>    <dbl>
1 (Intercept)    65.4      2.67     24.5 1.67e-34
2 OppFPI_PrevYear 0.445     0.242     1.84 7.08e- 2
```

```
glance(prev_fpi_glm)$adj.r.squared
```

```
[1] 0.03473927
```

The scatterplot above shows a fairly weak yet positive correlation between home-game attendance and the FPI of the opponent at the end of the previous season.

The linear model gives the slope of the linear fit depicted in the scatterplot. The model gives a slope of roughly 0.44497, which signifies that for every increase of 1 in the opponent's previous-season FPI, stadium attendance (as a percentage of Wallace Wade's total capacity) is predicted to increase by 0.44497% on average. The model indicates that this slope has a p-value of about 0.071, which is less than 0.1 and is significant given the difficulty of predicting future football attendance.

The adjusted r-squared value of about 0.0347 is very low, indicating that while a positive correlation is likely between attendance and opponent previous-season FPI, attendance is likely to also be based on other factors.

Previous-Season FPI Difference Between Duke & Opponent

```
prev_fpi_diff_glm <- linear_reg() |>
  set_engine("lm") |>
  fit(AttPct ~ OppFPI_PrevYear + FPI_Diff_PrevYear, data = home_att_data_prevFPI)

tidy(prev_fpi_diff_glm)
```

```
# A tibble: 3 x 5
  term          estimate std.error statistic  p.value
<chr>          <dbl>     <dbl>     <dbl>   <dbl>
1 (Intercept)    65.5       2.66      24.6 2.52e-34
2 OppFPI_PrevYear  0.843      0.401      2.10 3.94e- 2
3 FPI_Diff_PrevYear -0.455     0.366     -1.24 2.18e- 1
```

```
glance(prev_fpi_diff_glm)$adj.r.squared
```

```
[1] 0.0427744
```

When considering the FPI difference between Duke and its opponent in the previous season to a game, this factor is best represented in the model as an additive variable. The model gives a slope of roughly 0.843, which signifies that for every increase of 1 in the opponent's previous-season FPI, stadium attendance (as a percentage of Wallace Wade's total capacity) is predicted to increase by 0.843% on average. This is greater than the previous model, and this slope is also more significant ($p = 0.0394$).

Additionally, this model indicates that when the difference in previous-season FPI increases between Duke and its opponent increases (AKA when a matchup is more difficult for Duke based on the previous-season teams), stadium attendance decreases. However, the p-value for this is roughly 0.2183, suggesting that this trend may be due to chance rather than this association truly existing overall.

The adjusted r-squared value of this model is higher than the previous, suggesting that when you consider the FPI difference in addition to the opponent team's FPI, the model better predicts variation in stadium attendance.